

| H5: Optimization | |
|---|--|
| 5.1 introduction | |
| 5.1.1 concept and notation | |
| concept of optimization | <p>= problem of determining an argument for which a given function has an extreme value on a given domain</p> <p>formally: given a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, and a set $S \subseteq \mathbb{R}^n$, we seek $\mathbf{x}^* \in S$ such that f attains a minimum on S at \mathbf{x}^*, i.e. $f(\mathbf{x}^*) \leq f(\mathbf{x})$ for all $\mathbf{x} \in S$.</p> |
| objective function f | <p>= the function f for which we want to find an extremum</p> <p>> may be linear or non-linear, but is always differentiable</p> |
| feasible set S | <p>= set of equations and inequalities which put constraints on f</p> <p>> any \mathbf{x} which satisfies the constraints is called a feasible point, ie $\mathbf{x} \in S$</p> <p>If $S = \mathbb{R}^n$, the problem is unconstrained</p> |
| types of optimization problems | <p>A continuous opt. prob. has the form:</p> $\min_{\mathbf{x}} f(\mathbf{x}) \quad \text{subject to} \quad \mathbf{g}(\mathbf{x}) = \mathbf{0} \quad \text{and} \quad \mathbf{h}(\mathbf{x}) \leq \mathbf{0}$ <p>where $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ and $\mathbf{h} : \mathbb{R}^n \rightarrow \mathbb{R}^p$.</p> <p>is f, \mathbf{g} and \mathbf{h} are linear, the opt. problem is called linear programming if not, its nonlinear programming</p> |
| 5.2 optimality conditions | |
| 5.2.1 unconstrained optimality conditions | |
| first-order necessary condition | <p>We have a minimum in \mathbf{x}^* if:</p> $\nabla f(\mathbf{x}^*) = \mathbf{0}$ <p>> is a necessary condition, but doesn't prove \mathbf{x}^* is a minimum ie: it can be a maximum, minimum or saddle point</p> |
| second-order sufficient condition | <p>Consider the Hessian:</p> $\mathbf{H}_f(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f(\mathbf{x})}{\partial x_1^2} & \frac{\partial^2 f(\mathbf{x})}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f(\mathbf{x})}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f(\mathbf{x})}{\partial x_1 \partial x_2} & \frac{\partial^2 f(\mathbf{x})}{\partial x_2^2} & \cdots & \frac{\partial^2 f(\mathbf{x})}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f(\mathbf{x})}{\partial x_n \partial x_1} & \frac{\partial^2 f(\mathbf{x})}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f(\mathbf{x})}{\partial x_n^2} \end{bmatrix}$ <p>Now: At a critical point \mathbf{x}^*, where $\nabla f(\mathbf{x}) = \mathbf{0}$, if $\mathbf{H}_f(\mathbf{x}^*)$ is...</p> <ul style="list-style-type: none"> • Positive definite, then \mathbf{x}^* is a minimum of f • Negative definite, then \mathbf{x}^* is a maximum of f • Indefinite, then \mathbf{x}^* is a saddle point of f • Singular, then various pathological situations can occur |
| 5.3 optimization in one dimension | |
| unimodal function | <p>= a function $f: \mathbb{R} \rightarrow \mathbb{R}$ on interval $[a, b]$ for which there is a unique x^* AND for any x_1, x_2: $x_2 < x^*$ implies $f(x_1) > f(x_2)$ and $x_1 > x^*$ implies $f(x_1) < f(x_2)$</p> <p>ie: $f(x)$ is strictly increasing for $x < x^*$ and strictly increasing for $x > x^*$</p> |

| | |
|--|--|
| 5.3.1 golden section search | |
| golden section search | <p>For f unimodal on $[a,b]$ For $x_1, x_2 \in [a,b]$ with $x_1 < x_2$</p> <p>> compare the function values $f(x_1)$ and $f(x_2)$ > exclude either $[x_2, b]$ or $[a, x_1]$: - $f(x_1) < f(x_2) \rightarrow$ minimum on $[a, x_1]$ - $f(x_1) > f(x_2) \rightarrow$ minimum on $[x_2, b]$ > repeat this until you find the minimum</p> <p>To optimize this, each pair of points has the same relative position as the old pair > choose the this distance as τ = golden ratio = $(\sqrt{5} - 1)/2 \approx 0.618$ and $1-\tau$</p> <p>>> will converge if function is unimodal within initial bracket</p> |
| 5.3.2 successive parabolic search | |
| successive parabolic search | <p>evaluate the function at three points > fit a parabola to the points > use the minimum of the parabola as a new approximate value of the minimum</p> <p>>> algorithm isn't guaranteed to converge</p> |
| 5.3.3 Newton's method | |
| Newton's method | <p>Because of the Taylor expansion, the function is locally approximated by:</p> $f(x+h) \approx f(x) + f'(x)h + \frac{1}{2}f''(x)h^2$ <p>thus the minimum is given by:</p> $-f'(x)/f''(x),$ <p>use this to find the minimum of the objective function in an iterative way</p> <p>>> only converges if it started sufficiently close to the minimum > might converge to a maximum or inflection point</p> |
| 5.4 multidimensional unconstrained optimization | |
| 5.4.1 direct search | |
| direct search | <p>analogous to golden section search ie: objective function values are compared to one another > however isn't guaranteed to converge</p> |
| ex: Nelder and Mead | <p>for a function $f: \mathbb{R}^n \rightarrow \mathbb{R}$, choose $n+1$ points that form a simplex ie: $n+1$ points that aren't colinear > generate a new point along the straight line connecting the point with the highest function value and the centroid of the remaining n points > this new point replaces the other one</p> <p>>> repeat until converge</p> |
| 5.4.2 steepest descent | |
| steepest descent | <p>if we go in a direction $-\nabla f(\mathbf{x})$ we go into a descending direction > thus for a fixed point \mathbf{x} and direction $\mathbf{s} = -\nabla f(\mathbf{x})$, define</p> $\phi(\alpha) = f(\mathbf{x} + \alpha \mathbf{s})$ <p>> now we have a one-dimensional optimization problem > can be solved using previous methods</p> <p>>> repeat until converge</p> |

| | |
|--|---|
| 5.4.3 Newton's method | |
| Newton's method | <p>Approximate via Taylor:</p> $f(\mathbf{x} + \mathbf{s}) \approx f(\mathbf{x}) + \nabla f(\mathbf{x})^T \mathbf{s} + \frac{1}{2} \mathbf{s}^T \mathbf{H}_f(\mathbf{x}) \mathbf{s}, \quad (1)$ <p>where $\mathbf{H}_f(\mathbf{x})$ is the <i>Hessian matrix</i>. This quadratic function in \mathbf{s} is minimized when</p> $\mathbf{H}_f(\mathbf{x}) \mathbf{s} = -\nabla f(\mathbf{x}). \quad (1)$ <p>> unreliable method, unless started close enough to the solution</p> |
| 5.4.4 quasi-Newton methods | |
| quasi-Newton methods | <p>to lower calculations per iteration, consider using a general form:</p> $\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \mathbf{B}_k^{-1} \nabla f(\mathbf{x}_k)$ <p>with α_k a line search parameter \mathbf{B}_k some approximation of the Hessian matrix</p> |
| 5.4.5 secant updating methods | |
| secant updating methods | <p>methods which preserve symmetry in the approximate Hessian and positive definiteness > reduces work and storage by half + quasi-Newton step will be in descent >> most effective: BFGS (see course)</p> |
| 5.4.6 conjugate gradient method | |
| conjugate gradient method | <p>= alternative to Newton's method that doesn't use second derivatives > doesn't even use an approximation to Hessian matrix > suitable for large matrices</p> <p>uses gradient, but avoids repeatedly searching in the same directions > modifies the new gradient at each step to remove components in previous directions > sequence of conjugate: $(\mathbf{x}, \mathbf{y}) = \mathbf{x}^T \mathbf{H}_f \mathbf{y}$</p> |
| 5.5 nonlinear least squares | |
| nonlinear least squares problem | <p>view the problem as a special case of nonlinear optimization > given data points (t_i, y_i), $i=1, \dots, m$ we wish to find a vector $\mathbf{x} \in \mathbb{R}^n$ that best fits the model: $f(t, \mathbf{x})$, where $f: \mathbb{R}^{n+1} \rightarrow \mathbb{R}$.</p> <p>> define the components of the residual function $\mathbf{r}: \mathbb{R}^n \rightarrow \mathbb{R}^m$ by:</p> $r_i(\mathbf{x}) = y_i - f(t_i, \mathbf{x}), \quad i = 1, \dots, m,$ <p>then we wish to minimize the function</p> $\phi(\mathbf{x}) = \frac{1}{2} \mathbf{r}(\mathbf{x})^T \mathbf{r}(\mathbf{x})$ <p>Apply Newton's method: \mathbf{x}_k is an approximate solution > \mathbf{s}_k is given by the linear system:</p> $\mathbf{H}_\phi(\mathbf{x}_k) \mathbf{s}_k = -\nabla \phi(\mathbf{x}_k)$ <p>where the gradient vector and Hessian matrix of ϕ are given by</p> $\nabla \phi(\mathbf{x}) = \mathbf{J}^T(\mathbf{x}) \mathbf{r}(\mathbf{x})$ <p>and</p> $\mathbf{H}_\phi(\mathbf{x}) = \mathbf{J}^T(\mathbf{x}) \mathbf{J}(\mathbf{x}) + \sum_{i=1}^m r_i(\mathbf{x}) \mathbf{H}_{r_i}(\mathbf{x})$ <p>in which $\mathbf{J}^T(\mathbf{x})$ is the Jacobian matrix of $\mathbf{r}(\mathbf{x})$, and $\mathbf{H}_{r_i}(\mathbf{x})$ denotes the Hessian matrix of the component function $r_i(\mathbf{x})$.</p> <p>The Newton step \mathbf{s}_k is thus given by the linear system</p> $\left(\mathbf{J}^T(\mathbf{x}_k) \mathbf{J}(\mathbf{x}_k) + \sum_{i=1}^m r_i(\mathbf{x}_k) \mathbf{H}_{r_i}(\mathbf{x}_k) \right) \mathbf{s}_k = -\mathbf{J}^T(\mathbf{x}_k) \mathbf{r}(\mathbf{x}_k)$ |

| 5.5.1 Gauss-Newton method | |
|--|--|
| Gauss-Newton method | <p>note that \mathbf{H}_{ri} is multiplied by the corresponding residual component r_i > should be really small at a solution > drop this term:</p> $(\mathbf{J}^T(\mathbf{x}_k)\mathbf{J}(\mathbf{x}_k))\mathbf{s}_k = -\mathbf{J}^T(\mathbf{x}_k)\mathbf{r}(\mathbf{x}_k)$ <p>determines an approximate Newton step \mathbf{s}_k at each iteration.</p> <p>Recognize this system as the normal equations for the $m \times n$ linear least squares problem:</p> $\mathbf{J}(\mathbf{x}_k)\mathbf{s}_k \cong -\mathbf{r}(\mathbf{x}_k)$ <p>which can be solved more reliably by orthogonal factorization of $\mathbf{J}(\mathbf{x}_k)$ > next approximate solution is given by: $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{s}_k$ > repeat until convergence</p> |
| 5.5.2 Levenberg-Marquardt method | |
| Levenberg-Marquardt method | <p>at each iteration, the linear system for the step \mathbf{s}_k is of the form:</p> $(\mathbf{J}^T(\mathbf{x}_k)\mathbf{J}(\mathbf{x}_k) + \mu_k\mathbf{I})\mathbf{s}_k = -\mathbf{J}^T(\mathbf{x}_k)\mathbf{r}(\mathbf{x}_k)$ <p>with μ_k a nonnegative scalar parameter chosen by some strategy > corresponding linear least squares problem is:</p> $\begin{bmatrix} \mathbf{J}(\mathbf{x}_k) \\ \sqrt{\mu_k}\mathbf{I} \end{bmatrix} \mathbf{s}_k \cong \begin{bmatrix} -\mathbf{r}(\mathbf{x}_k) \\ \mathbf{0} \end{bmatrix}$ |
| 5.6 constrained optimization | |
| constrained optimization problem | <p>= optimization problem for which the minimum is located outside the feasible set > solution occurs on the boundary of the feasible set</p> <p>They can be linear and nonlinear and subdivided into two categories:</p> <ul style="list-style-type: none"> • Equality constraints which are of the general form $\mathbf{g}(\mathbf{x}) = \mathbf{0}$ • Inequality constraints which are of the general form $\mathbf{h}(\mathbf{x}) \leq \mathbf{0}$ <p>>> can be solved using Lagrange multipliers</p> |
| 5.6.1 the trust-region constrained algorithm | |
| the trust-region constrained algorithm | <p>The minimize function of <i>scipy.optimize</i> provides several algorithms for constrained min > one of those is the t-r constr. alg. which deals with constr. min. problems of the form:</p> $\begin{aligned} &\min_{\mathbf{x}} f(\mathbf{x}) \\ &\text{subject to: } \mathbf{c}_l \leq \mathbf{c}(\mathbf{x}) \leq \mathbf{c}_u \\ &\quad \mathbf{x}_l \leq \mathbf{x} \leq \mathbf{x}_u \end{aligned}$ <p>When $c_{l,j} = c_{u,j}$ the method reads the j^{th} constraint as an equality constraint and deals with it accordingly. Beside that, one-sided constraints can be specified by setting the upper (<i>u</i>) or lower (<i>l</i>) bound to <code>np.inf</code> with the appropriate sign. As an illustration this method will be applied to the Rosenbrock example.</p> |